Chronikis User Manual[*]

Kevin S. Van Horn
Adobe Inc.

March 29, 2019

Chronikis (kroh-NEE-kees) is a special-purpose language for creating time-series models. It comes with a compiler `chronikisc`, and an R package `chronikis` that contains utilities for compiling Chronikis programs as well as estimating and forecasting with the compiled time-series models.

The name "Chronikis" is derived from the phrase χρονική σειρά (*chronikí seirá*), which means "time series" in Greek.

This initial release is still missing a number of functions and distributions that the language ought to have; the focus was on implementing enough that all of the models in the `compiler/Acceptance` *subdirectory could be compiled.*

1 Installation

1. Install needed prerequisites.

   (a) Install R if necessary.

   (b) Install the R package rstan.

   (c) Install The Haskell Tool Stack.

2. Clone or download this git repository and `cd` to the root.

3. `cd compiler`

4. `stack install`
   You may run into issues with the Haskell package `hmatrix-gsl`. This requires the GNU Scientific Library and the `pkg-config` utility. On Mac OSX with Homebrew you can obtain these by issuing the commands

   ```
   brew install gsl
   brew install pkg-config
   ```

5. Fire up R and install the package `chronikis` using these commands:
   ```
   setwd(path_to_cloned_git_repository)
   install.packages('chronikis_0.1.0.tar.gz', repos=NULL)
   ```

---

[*]©2019, Adobe Inc. This document is licensed to you under the Apache License, Version 2.0. You may obtain a copy of the license at http://www.apache.org/licenses/LICENSE-2.0.

6. Add `~/.local/bin` to your bash `PATH` variable.

7. Add `~/.local/bin` to your R `PATH` variable. (This is where `chronikisc` got installed.) To do this, open your `~/.Renviron` file (create it if it doesn't exist) and add the line
`PATH=${PATH}:~/.local/bin`
If there is already a line setting `PATH`, add `~/.local/bin` to the end.

2   Using Chronikis

Here is what the process of compiling and running a model generally looks like. In the following, `ytrain` is the time series on which to train the model. The following assumes that you have written a Chronikis program and saved it to "`my_model.cks`".

```
library(rstan)
library(chronikis)
sm <- cksCompile("my_model.cks", "createSSMs")
source("my_model.R")  # file was created by cksCompile

# Model training
margs <- mdlArgs(model, arguments, go, here, ...)
sma <- setArgs(sm, margs)
fit <- hmc_estimate(ytrain, sma)
# fit <- vb_estimate(ytrain, sma)
# fit <- map_estimate(ytrain, sma)
post <- posterior_sample(fit)
npost <- 100
# It gets slow if you use more than 100 posterior draws
models0 <- createSSMs(margs, post, npost)
filtered0 <- filter_models(ytrain, models0)

# Model checking
lltrain <- averagedLL(models0, ytrain)
# Compare yrep[,i], 1 <= i <= npost, to ytrain;
# are they visually similar?
# Can also compute various summary statistics T()
# and check that T(ytrain) is within range of values
# T(yrep[,i]).
yrep <- forecast_sample(models0, length(ytrain), 1)
# Check that smoothed time series looks reasonable.
s <- smoothed_ts(ytrain, filtered0)
ysmooth <- average_normals(s$means, s$stddevs, 0.1, TRUE)
# Get one-step predictive residuals.
res <- lapply(filtered, residuals, type='raw', sd=FALSE)
```

```
# Forecasting
models <- update_models(filtered = filtered0)
nsteps <- 20
alpha <- 0.10
# 90% predictive intervals
fc <- forecast_intervals(models, nsteps, alpha, TRUE)
# fc$mean[t] is the predictive mean t steps in the future
# (fc$lower[t], fc$upper[t]) is the 90% predictive interval
# t steps in the future.
ndpm <- 200  # Number of predictive draws per model
ymc <- forecast_sample(models, nsteps, ndpm)
# ymc is a nsteps x (npost * ndpm) matrix of predictive draws
# that can be used for Monte Carlo analysis
minmc <- matrixStatts::colMins(ymc[11:20, ])
# minmc is a posterior predictive sample of
# min(y[n+11], ..., y[n+20]), where n = length(ytrain)
pred <- quantile(minmc, c(0.05, 0.95))
# pred is 90% predictive interval for
# min(y[n+11], ..., y[n+20])
```

Use `help(package=chronikis)` to see documentation on all of the functions provided in the R package. The directory `compiler/Acceptance` contains examples of Chronikis programs.

3 Language Definition

3.1 Overview

A Chronikis program defines a parameterized distribution over time series (infinite sequences of real numeric values.) The general form is

$$\text{def main (} knownParameters \text{) = } TSDExpression$$

where *knownParameters* is a comma-separated list of variables with their types, and *TSDExpression* is an expression that denotes a probability distribution over time series.

Here are some examples for *knownParameters*:

- `N: int{1,}, mu: real[N], sigma: real{0.0,}`
  This declares `N` to be a positive integer, `mu` to be a real-valued vector of length N, and `sigma` to be a nonnegative real value.

- `rho: real{0.0, 1.0}, sigma_a, sigma_h: real{0.0,}[3]`
  This declares `rho` to be a real number between 0 and 1, and `sigma_h` and `sigma_a` to be length-3 vectors of nonnegative real values.

`N`, `mu`, etc. are called *known* parameters because they are already known when the model is created, rather than being inferred from training data. The range

constraints (e.g. `{0.0,}`, `{0.0,1.0}`) on known parameters are checked when a model is trained.

A *TSDExpression* can have one of three forms:

- *variable = defExpr ; TSDExpression*$_1$
  This defines the *variable* to have the value *defExpr* within *TSDExpression*$_1$.

- *variable ~ distrExpression ; TSDExpression*$_1$
  This defines the *variable* as a latent (unobserved) variable having the probability distribution *distrExpression*, with the time-series distribution *TSDExprresion*$_1$ being conditional on the value of the *variable*. Typically *variable* will be a *fitted* parameter, i.e. a parameter to be inferred from training data, with *distrExpr* being the prior distribution for *variable*.

- A function call (including use of the binary operator `+`) that returns a distribution over time series. Some examples:

  - `wn(sigma)` is a white noise process with mean 0 and variance `sigma` squared. That is, it corresponds to independent normal distributions for each time step.

  - `qp(7.0, ell, 6, 0.0, sigma_p) + constp(mu0, sigma0)` is a distribution over periodic patterns of period 7, with the distribution having smoothness parameter `ell`, and scale parameter `sigma_p`, and the repeating pattern having an unknown mean that is given a `normal(mu0,sigma0)` prior.

The above is a recursive definition, so in practice a Chronikis program will have one or more variable definitions and/or variable draws, followed by a function call that returns a time-series distribution.

3.2 Examples

(To be written. For now, look in the directory `compiler/Acceptance`.)

3.3 Types

Chronikis programs are strongly and statically typed, but the only place you provide types are in the parameter list of `main`. The compiler infers all the remaining types.

A type has four attributes:

- its *element type*, which is `int` or `real`;

- its *shape*, which is a list of nonnegative integers;

- whether or not it is a *time series*, indicated with `$`; and

- whether or not it is a *distribution*, indicated with `~`.

4

The attributes apply in the order given above; thus,

- "distribution over time series of real vectors of length n" is a valid type (`real[n]$~`), but

- "time series of distributions over real vectors of length n" (`real[n]~$`) is *not* valid, and

- "length-n vector of distributions over univariate time series" (`real$~[n]`) is *not* valid.

Some examples:

- `real` is the type of real scalars.

- `real[n]` is the type of vectors of length n.

- `real[m,n]` is the type of m-by-n matrices.

- `real~` is the type of distributions over `real` values; an example is `normal(mu, sigma)`.

- `real[k,p]$` is the type of k-by-p matrix-valued time series.

- `real$~` is the type of distributions over univariate time series.

Currently, variables and parameters cannot have types that are time series or distributions, but we can construct expressions having such types.

3.4 Data Functions

In the following, a *signature* gives argument types in parentheses, and the corresponding return type after a colon. Italicized variables such as m or n indicate that there is one such signature for every combination of nonnegative values of the variables.

- (+) (binary operator). Signatures:

    - (`int`, `int`): `int`.
    - (`real`, `real`): `real`.
    - (`real`, `real[`n`]`): `real[`n`]`.
      Add a scalar to each element of a vector.
    - (`real[`n`]`, `real`): `real[`n`]`.
      Add a scalar to each element of a vector.
    - (`real`, `real[`m`,`n`]`): `real[`m`,`n`]`.
      Add a scalar to each element of a matrix.
    - (`real[`m`,`n`]`, `real`): `real[`m`,`n`]`.
      Add a scalar to each element of a matrix.

- (-) (unary operator). Negation. Signatures:

    - (int): int.
    - (real): real.
    - (real[n]): real[n].
      Negate each element of a vector.
    - (real[m,n]): real[m,n].
      Negate each element of a matrix.

- (-) (binary operator). Signatures:

    - (int, int): int.
    - (real, real): real.
    - (real, real[n]): real[n].
      Subtract each element of a vector from a scalar.
    - (real[n], real): real[n].
      Subtract a scalar from each element of a vector.
    - (real, real[m,n]): real[m,n].
      Subtract each element of a matrix from a scalar.
    - (real[m,n], real): real[m,n].
      Subtract a scalar from each element of a matrix.

- (*) (binary operator). Signatures:

    - (int, int): int.
    - (real, real): real.
    - (real, real[n]): real[n].
      Multiply each element of a vector by a scalar.
    - (real[n], real): real[n].
      Multiply each element of a vector by a scalar.
    - (real, real[m,n]): real[m,n].
      Multiply each element of a matrix by a scalar.
    - (real[m,n], real): real[m,n].
      Multiply each element of a matrix by a scalar.
    - (real, real[m,n,p]): real[m,n,p].
      Multiply each element of a 3D array by a scalar.
    - (real[m,n,p], real): real[m,n,p].
      Multiply each element of a 3D array by a scalar.

- (/) (binary operator). Signatures:

    - (real, real): real.
    - (real, real[n]): real[n].
      Divide a scalar by each element of a vector.

- $(\texttt{real}[n], \texttt{real})$: $\texttt{real}[n]$.
  Divide each element of a vector by a scalar.

- $(\texttt{real}, \texttt{real}[m,n])$: $\texttt{real}[m,n]$.
  Divide a scalar by each element of a matrix.

- $(\texttt{real}[m,n], \texttt{real})$: $\texttt{real}[m,n]$.
  Divide each element of a matrix by a scalar.

- (^) (binary operator). Exponentiation. Signatures:

  - $(\texttt{real}, \texttt{int})$: $\texttt{real}$.

- ([]) Array indexing, 1-based. Signatures:

  - $(\texttt{real}[n], \texttt{int})$: $\texttt{real}$.
    $a[i]$ is element i of vector a.

  - $(\texttt{real}[m,n], \texttt{int})$: $\texttt{real}[n]$.
    $a[i]$ is row i of matrix a (as a vector).

  - $(\texttt{real}[m,n], \texttt{int}, \texttt{int})$: $\texttt{real}$.
    $a[i,j]$ is row i, column j of matrix a.

  - etc. for higher-dimensional arrays.

- - ({}) Array enumeration. Signatures:

  - $(\texttt{real}[m,n], \ldots, \texttt{real}[m,n])$: $\texttt{real}[k,m,n]$.
    $\{M_1, \ldots, M_k\}$ is the 3D array a such that $a[i] = M_i$ for $1 \leq i \leq k$.
    It is required that $k \geq 1$.

  - $(\texttt{real}[m,n,p], \ldots, \texttt{real}[m,n,p])$: $\texttt{real}[k,m,n,p]$.
    $\{A_1, \ldots, A_k\}$ is the 4D array a such that $a[i] = A_i$ for $1 \leq i \leq k$. It
    is required that $k \geq 1$.

  - etc. for higher dimensions.

- $\texttt{blocks4}$. Create a matrix from four submatrices. Signatures:

  - $(\texttt{real}[m_1,n_1], \texttt{real}[m_1,n_2], \texttt{real}[m_2,n_1], \texttt{real}[m_2,n_2])$: $\texttt{real}[m,n]$
    where $m = m_1 + m_2$ and $n = n_1 + n_2$.
    $\texttt{blocks4}(A, B, C, D)$ is the matrix

    $$\begin{pmatrix} A & B \\ C & D \end{pmatrix}.$$

  - $(\texttt{real}, \texttt{real}[n], \texttt{real}[m], \texttt{real}[m,n])$: $\texttt{real}[m+1,n+1]$.
    $\texttt{blocks4}(a, b, c, D) = \texttt{blocks4}(A, B, C, D)$ where

    * A is the 1×1 matrix created from scalar a,
    * B is the $1 \times n$ matrix created from vector b, and
    * C is the $m \times 1$ matrix created from vector c.

– $(\texttt{real}[m,n], \texttt{real}[m], \texttt{real}[n], \texttt{real})$: $\texttt{real}[m+1,n+1]$.
    $\texttt{blocks4}(A, b, c, d) = \texttt{blocks4}(A, B, C, D)$ where

      * B is the $m \times 1$ matrix created from vector b,
      * C is the $1 \times n$ matrix created from vector c,
      * D is the 1×1 matrix created from scalar d.

- `diag`. Create a diagonal or block-diagonal matrix. Signatures:

  – $()$: $\texttt{real}[0,0]$.
    $\texttt{diag}()$ is the 0×0 matrix.
  – $(\texttt{real}, \ldots, \texttt{real})$: $\texttt{real}[k,k]$.
    $\texttt{diag}(x_1, \ldots, x_k) = \texttt{diag}(\texttt{vec}(x_1, \ldots, x_k))$ for $k \geq 1$.
  – $(\texttt{real}[n])$: $\texttt{real}[n,n]$.
    $\texttt{diag}(v)$ is the diagonal matrix with v as its diagonal.
  – $(\texttt{real}[m_1,n_1], \ldots, \texttt{real}[m_k,n_k])$: $\texttt{real}[m,n]$
    where $m = m_1 + \cdots + m_k$ and $n = n_1 + \cdots + n_k$ and $k \geq 1$.
    $\texttt{diag}(M_1, \ldots, M_k)$ is the block diagonal matrix having $M_1, \ldots, M_k$
    as the blocks.
  – $(\texttt{real}[k,m,n])$: $\texttt{real}[km,kn]$.
    $\texttt{diag}(A) = \texttt{diag}(A[1], \ldots, A[k])$.
  – $(T_1, \ldots, T_k)$: $\texttt{real}[m,n]$
    where each T_i is one of $\texttt{real}$, $\texttt{real}[n_i]$, $\texttt{real}[m_i,n_i]$, or $\texttt{real}[p_i,m_i,n_i]$,
    and m and n are computed from the shapes of the types T_i.
    $\texttt{diag}(a_1, \ldots, a_k) = \texttt{diag}(M_1, \ldots, M_k)$ where M_i is

      * $\texttt{diag}(a_i)$ if T_i is $\texttt{real}$ or $\texttt{real}[n_i]$ or $\texttt{real}[p_i,m_i,n_i]$;
      * a_i if T_i is $\texttt{real}[m_i,n_i]$.

- `diag_sqr`. Signatures: same as `diag`.
  $\texttt{diag\_sqr}(x_1, \ldots, x_k) = \texttt{diag}(\texttt{square}(x_1), \ldots, \texttt{square}(x_k))$.

- `div` (binary operator). Integer division. Signatures:

  – $(\texttt{int}, \texttt{int})$: $\texttt{int}$.
    It is required that $b \neq 0$ in the expression $a \,\texttt{div}\, b$. Rounds down for
    nonnegative arguments.

- `i2r`. Convert an `int` to a `real`. Signatures:

  – $(\texttt{int})$: $\texttt{real}$.

- `mat11`. Create a 1×1 matrix. Signatures:

  – $(\texttt{real})$: $\texttt{real}[1,1]$.

- `mat22`. Create a 2×2 matrix. Signatures:

– $(\texttt{real}, \texttt{real}, \texttt{real}, \texttt{real})$: $\texttt{real[2,2]}$.
  $\texttt{mat22}(a, b, c, d)$ is the matrix
  $$\left(\begin{array}{cc} a & b \\ c & d \end{array} \right).$$

- `negate`. Same as unary (-).

- `sqrt`. Square-root applied element-wise. Signatures:

  – $(\texttt{real})$: $\texttt{real}$.
  – $(\texttt{real}[n])$: $\texttt{real}[n]$.
  – $(\texttt{real}[m,m])$: $\texttt{real}[m,n]$.
  – etc. for higher dimensions.

- `square`. Square each element. Signatures:

  – $(\texttt{real})$: $\texttt{real}$.
  – $(\texttt{real}[n])$: $\texttt{real}[n]$.
  – $(\texttt{real}[m,n])$: $\texttt{real}[m,n]$.
  – etc. for higher dimensions.

- `to_matrix`. Convert a vector to a matrix. Signatures:

  – $(\texttt{real}[n])$: $\texttt{real}[n,\texttt{1}]$.

- `transp`. Matrix transpose. Signatures:

  – $(\texttt{real}[m,n])$: $\texttt{real}[n,m]$.

- `vec`. Create a vector. Signatures:

  – $(\texttt{real}, \ldots, \texttt{real})$: $\texttt{real}[k]$.
  $\texttt{vec}(x_1, \ldots, x_k)$ is the vector of length $k \geq 0$ with the indicated elements.
  – $(\texttt{real}[n_1], \ldots, \texttt{real}[n_k])$: $\texttt{real}[n]$
  where $n = n_1 + \cdots + n_k$ and $k \geq 1$.
  $\texttt{vec}(v_1, \ldots, v_k)$ is the result of appending vectors v_1 through v_k.
  – $(T_1, \ldots, T_k)$: $\texttt{real}[n]$
  where each T_i is $\texttt{real}$ or $\texttt{real}[n_i]$, and n is computed from the shapes of the types T_i.
  $\texttt{vec}(a_1, \ldots, a_k) = \texttt{vec}(v_1, \ldots, v_k)$ where v_i is

    * the length-one vector $\texttt{vec}(a_i)$ if T_i is $\texttt{real}$,
    * a_i if T_i is $\texttt{real}[n_i]$.

- `vec0`. Create a vector of zeroes. Signatures:

  – $(\texttt{int})$: $\texttt{real}[n]$
  where n is the argument.
  $\texttt{vec0}(n)$ is the length-n vector $\texttt{vec}(0, \ldots, 0)$.

3.5 Data Distributions

- `certainly`. Degenerate distribution that assigns probability 1 to its argument. An example usage would be
  `x ~ uniform(0.0,1.0);`
  `y ~ uniform(0.0, 2.0);`
  `certainly(x+y)`
  which is the distribution of the sum of draws from a $U(0,1)$ and a $U(0,2)$ distribution. Signatures:

  - (`real`): `real`˜.
  - (`real[`n`]`): `real[`n`]`˜.
  - (`real[`m,n`]`): `real[`m,n`]`˜.
  - etc. for higher dimensions.

- `exponential_m`. Exponential distribution parameterized by mean. Signatures:

  - (`real`): `real`˜.
    `exponential_m(`μ`)` is the exponential distribution with mean μ, which must be positive.

- `exponential_mt`. Truncated exponential distribution parameterized by mean and upper bound. Signatures:

  - (`real`,`real`): `real`˜.
    `exponential_mt(`μ, u`)` is the truncated exponential distribution with upper bound u and mean μ. It is required that $u > \mu > 0$. Note that `exponential_mt(`μ, u`)` is *not* the same as taking `exponential_m(`μ`)` and truncating above at u, as the mean of that distribution is less than μ.

- `exponential_r`. Exponential distribution parameterized by rate. Signatures:

  - (`real`): `real`˜.
    `exponential_r(`θ`)` = `exponential_m(`$1/\theta$`)`. It is required that $\theta > 0$.

- `exponential_rt`. Truncated exponential distribution parameterized by rate and upper bound. Signatures:

  - (`real`,`real`): `real`˜.
    `exponential_rt(`θ, u`)` is the distribution `exponential_r(`θ`)` truncated above at u. Note that its mean is *not* $1/\theta$, due to the truncation. It is required that $\theta > 0$ and $u > 0$.

- `half_cauchy`. Cauchy distribution truncated below at 0. Signatures:

– (`real`): `real~`.
  `half_cauchy(s)` is the Cauchy distribution with scale parameter $s > 0$, truncated below at 0.

- `half_normal`. 0-mean normal distribution truncated below at 0. Signatures:

  – (`real`): `real~`.
    `half_normal(`σ`)` is the normal distribution with mean 0 and standard deviation $\sigma > 0$, truncated below at 0.

- `normal`. Univariate normal distribution. Signatures:

  – (`real, real`): `real~`.
    `normal(`μ, σ`)` is the normal (Gaussian) distribution with mean μ and standard deviation $\sigma > 0$.

- `uniform`. Uniform distribution. Signatures:

  – (`real, real`): `real~`.
    `uniform(`l, u`)` is the uniform distribution over the interval from l to u. It is required that $l < u$.

3.6   Time-series Distributions

In the following we write y_t for the value of a time series at time t. Note that time series start at $t = 1$, so when we refer to y_0 below, this is a latent "one step before the first" value.

- (+) (binary op). Sum of time-series distributions. Signatures:

  – (`real$~, real$~`): `real$~`.
    The distribution $d_1 + d_2$ is equivalent to

    $$v_1 \sim d_1$$
    $$v_2 \sim d_2$$
    $$y_t = v_{1t} + v_{2t} \quad \text{for all } t \geq 1$$

- `accum`. Accumulate. Signatures:

  – (`real$~, real, real`): `real$~`.
    The distribution `accum(`d, μ, σ`)` is eqivalent to

    $$\delta \sim d$$
    $$y_0 \sim \texttt{normal}(\mu, \sigma)$$
    $$y_t = y_{t-1} + \delta_t \quad \text{for all } t \geq 1$$

  Put another way, the time series of differences $y_t - y_{t-1}$ has the distribution d. It is required that $\sigma > 0$.

11

- `ar1`. An AR(1) process. Signatures:

  - $(\texttt{real}, \texttt{real}, \texttt{real})\colon \texttt{real\$\textasciitilde}$.
    $\texttt{ar1}(\phi, \sigma_q, \sigma_0)$ is equivalent to

    $$y_0 \sim \texttt{normal}(0, \sigma_0)$$
    $$y_t \sim \texttt{normal}(\phi y_{t-1}, \sigma_q) \quad \text{for all } t \geq 1$$

    If $\sigma_0^2 = \sigma_q^2 / \left(1 - \phi^2\right)$ then this distribution is a stationary process with mean 0, standard deviation σ_0, and correlation ϕ^k between y_t and y_{t+k}. It is required that $0 < \phi < 1$, $\sigma_q > 0$, and $\sigma_0 > 0$.

- `constp`. Constant time series. Signatures:

  - $(\texttt{real}, \texttt{real})\colon \texttt{real\$\textasciitilde}$.
    $\texttt{constp}(\mu, \sigma)$ is equivalent to

    $$y_0 \sim \texttt{normal}(\mu, \sigma)$$
    $$y_t = y_0 \quad \text{for all } t \geq 1$$

    It is required that $\sigma > 0$.

- `qp`. Quasi-periodic process. Signatures:

  - $(\texttt{real}, \texttt{real}, \texttt{int}, \texttt{real}, \texttt{real})\colon \texttt{real\$\textasciitilde}$.
    $\texttt{qp}(P, \ell, n, \rho, \sigma)$ is a distribution over quasiperiodic time series. The arguments P, ℓ, and n must currently be numeric literals, e.g. $\texttt{qp}(7.0, 0.8, 6, \rho, \sigma)$. It is required that $P > 0$, $\ell > 0$, $0 \leq n < P$, $0 \leq \rho \leq 1$, and $\sigma > 0$.

    * The marginal distribution for y_t is $\texttt{normal}(0, \sigma)$. Additionally, the periodic pattern itself is centered around 0.
    * P is the period; it does not need to be an integer.
    * ℓ gives a smoothness length scale; y_t and y_{t+k} are positively correlated for $|k|$ less than about $\ell P / 4$.
    * Nonzero values for ρ allow the periodic pattern to change somewhat from one period to another. Specifically, if P is an integer then y_t and y_{t+P} have a correlation coefficient of ϕ^P, where $\phi = \sqrt{1 - \rho^2}$.
    * d is the minimum required degrees of freedom for the periodic pattern. This is an implementation wart that will be removed in a later release of Chronikis; for now you can just set it to the minimum of 10 and $P - 1$.
    * This time-series distribution closely approximates a one-dimensional Gaussian process with covariance function

    $$\kappa(x) = \sigma^2 \phi^x \exp\left(-2\left(\frac{\sin(\pi x)}{\ell}\right)^2\right).$$

- `rw`. Random-walk process. Signatures:

  - $(\texttt{real}, \texttt{real}, \texttt{real})$: `real$~`.
    $\texttt{rw}(\mu_0, \sigma_0, \sigma_q)$ is equivalent to

    $$y_0 \sim \texttt{normal}(\mu_0, \sigma_0)$$
    $$y_t \sim \texttt{normal}(y_{t-1}, \sigma_q) \quad \text{for all } t \geq 1$$

    It is required that $\sigma_0 > 0$ and $\sigma_q > 0$.

- `ssm`. General form for creating a linear state-space model. Signatures:

  - $(\texttt{real}[m], \texttt{real}, \texttt{real}[m,m], \texttt{real}[m,m], \texttt{real}[m], \texttt{real}[m,m])$: `real$~`.
    $\texttt{ssm}(z, h, T, Q, a_0, P_0)$ specifies the following linear state-space model:

    $$\alpha_0 \sim \texttt{mvnormal}(a_0, P_0)$$
    $$\alpha_t \sim \texttt{mvnormal}(\alpha_{t-1}, Q) \quad \text{for all } t \geq 1$$
    $$y_t \sim \texttt{normal}(z'\alpha_t, \sigma) \quad \text{for all } t \geq 1$$
    $$\sigma = \sqrt{h}$$

    where $\texttt{mvnormal}(\mu, \Sigma)$ is the multivariate normal distribution with mean vector μ and covariance matrix Σ. It is required that $h > 0$ and T, Q, and P be nonnegative definite matrices.

- `wn`. White noise. Signatures:

  - $(\texttt{real})$: `real$~`.
    $\texttt{wn}(\sigma)$ is equivalent to

    $$y_t \sim \texttt{normal}(0, \sigma) \quad \text{for all } t \geq 1$$

    It is required that $\sigma > 0$.

3.7   Formal Syntax

In the following,

- "$*$" means *zero* or more instances separated by commas,

- "$+$" means *one* or more instances separated by commas,

- "?" means optional (zero or one instance),

- "|" separates alternatives, and

- text in typewriter font is literal text.

Here is the grammar:

$$
\begin{aligned}
Program &\coloneqq \texttt{def main (} \; Parms \; \texttt{) = } TSDExpr \\
TSDExpr &\coloneqq Expr \\
&\qquad (\textit{type must be } \texttt{real\$\textasciitilde}) \\
Parms &\coloneqq ParmGroup* \\
ParmGroup &\coloneqq Identifier+ : Type \\
Expr &\coloneqq DefExpr \mid DrawExpr \mid OpExpr \\
DefExpr &\coloneqq Variable \texttt{=} OpExpr \texttt{;} Expr \\
DrawExpr &\coloneqq Variable \texttt{\textasciitilde} OpExpr \texttt{;} Expr \\
&\qquad (\textit{the OpExpr must have a distribution type}) \\
OpExpr &\coloneqq Term \mid UnOp \; OpExpr \mid OpExpr \; BinOp \; OpExpr \mid OpExpr \; Index \\
UnOp &\coloneqq \texttt{+} \mid \texttt{-} \\
BinOp &\coloneqq \texttt{+} \mid \texttt{-} \mid \texttt{*} \mid \texttt{/} \mid \texttt{div} \mid \texttt{\textasciicircum} \\
Index &\coloneqq \texttt{[} \; IExpr+ \; \texttt{]} \\
IExpr &\coloneqq Expr \\
&\qquad (\textit{type must be } \texttt{int}) \\
Term &\coloneqq Variable \mid Literal \mid \texttt{(} \; Expr \; \texttt{)} \mid FctApp \mid Array \\
FctApp &\coloneqq FctName \; \texttt{(} \; Expr* \; \texttt{)} \\
Array &\coloneqq \texttt{\{} \; Expr+ \; \texttt{\}} \\
&\qquad (\textit{all entries must have the same type}) \\
Literal &\coloneqq IntegerLit \mid RealLit \\
Type &\coloneqq ElemType \; Bounds? \; Shape? \\
ElemType &\coloneqq \texttt{int} \mid \texttt{real} \\
Bounds &\coloneqq \texttt{\{} \; OpExpr? \; \texttt{,} \; OpExpr? \; \texttt{\}} \\
Shape &\coloneqq \texttt{[} \; OpExpr+ \; \texttt{]}
\end{aligned}
$$

Operator precedence, from lowest to highest, is

- binary +, binary - (left associative);

- *, /, div (left associative);

- ^ (right associative);

- unary +, unary -;

- indexing.

The following are tokens, so there is no whitespace within them:

$$
\begin{aligned}
\textit{IntegerLit} &\coloneqq \textit{Digits} \\
\textit{RealLit} &\coloneqq \textit{Digits Fraction? Exponent?} \\
&\qquad (\textit{must have at least one of Fraction or Exponent}) \\
\textit{Fraction} &\coloneqq \texttt{.}\ \textit{Digits} \\
\textit{Exponent} &\coloneqq \textit{ExpChar Sign? Digits} \\
\textit{ExpChar} &\coloneqq \texttt{E}\ |\ \texttt{e} \\
\textit{Sign} &\coloneqq \texttt{+}\ |\ \texttt{-} \\
\textit{Digits} &\coloneqq \textit{Digit}+ \\
\textit{Digit} &\coloneqq \texttt{0}\ |\ \cdots\ |\ \texttt{9} \\
\textit{Variable} &\coloneqq \textit{Identifier} \\
\textit{FctName} &\coloneqq \textit{Identifier} \\
\textit{Identifier} &\coloneqq \textit{IdentChar0 IdentChar}* \\
\textit{IdentChar0} &\coloneqq \texttt{a}\ |\ \cdots\ |\ \texttt{z}\ |\ \texttt{A}\ |\ \cdots\ |\ \texttt{Z} \\
\textit{IdentChar} &\coloneqq \textit{IdentChar0}\ |\ \textit{Digit}\ |\ \texttt{\_}
\end{aligned}
$$

Note that there is no automatic promotion of integers to reals.